

# AllelePerturb reveals measurement limits of protein-coding variant prediction in single-cell transcriptomics

Bo Li<sup>1,\*</sup>

<sup>1</sup>Department of Biomedical Engineering, University of Florida, Gainesville, FL, USA.

\*Correspondence: [bo.li@ufl.edu](mailto:bo.li@ufl.edu)

## Abstract

Single-cell perturbation prediction is becoming a central task for virtual-cell modeling, but current methods and benchmarks largely define perturbations at gene, drug or condition resolution. This design is suited to knockout and drug-response screens, yet does not address how different protein-coding variants within the same gene alter cellular state. Here we introduce AllelePerturb, a benchmark and evaluation framework for held-out protein-coding variant prediction in single-cell transcriptomics. AllelePerturb integrates 470 protein-coding variants across TP53, KRAS, GATA1 and JAK1, comprising 321,043 single cells (including wild-type controls) from Perturb-seq, base editing and scSNV-seq, and evaluates prediction quality across direction recovery, allele-level discrimination and differential-expression fidelity.

Across a family of 20 predictors (18 feature-model combinations spanning six regression heads and three feature spaces, plus two references), all feature-model predictors recover perturbation direction (Pearson- $\delta$  0.55–0.65), yet none reliably distinguishes held-out variants (PDS 0.49–0.52, statistically indistinguishable from chance). Four published perturbation models that accept a continuous variant representation (scGen, scVIDR, Biolord and CellFlow), scored through the identical harness, likewise sit at chance, whereas gene-keyed foundation models (GEARS, STATE and scGPT) are allele-blind by construction; none of the evaluated model classes exceeded the empirical discrimination floor under the available measurement regimes. Split-half analysis traces this direction-ranking dissociation to a measurement-resolution floor governed by variant-to-wild-type effect size, within-variant sampling noise and effective cells per variant. An oracle analysis separates measurement from model limits: for three genes even a perfect predictor cannot rank held-out variants at the available depth, whereas for JAK1 the measurement suffices yet models still fail; below this window a benchmark cannot recover the true ordering of models. Direction recovery is entirely gene-shared: decomposing each effect into a gene mean and an allele-specific residual leaves near-zero residual signal for every predictor, and modern mutation-aware and structure-aware representations, including ESM2 and AlphaMissense, do not overcome the floor. The same floor appears in gene-level Perturb-seq atlases, where roughly half of Replogle perturbations are un-rankable under the specified split-half energy-distance criterion at native depth. AllelePerturb reframes variant-level perturbation prediction as a power-aware evaluation problem and provides effect-size-conditioned triage and diagnostic guidance for future variant-level screens.

---

## Introduction

Single-cell perturbation prediction is becoming a central testbed for virtual-cell modeling. High-throughput perturbation screens now make it possible to measure how cell populations respond to genetic and chemical interventions at scale, and computational models are increasingly expected to predict the

transcriptional consequences of perturbations that have not been experimentally observed. A growing ecosystem of generative, optimal-transport, graph-based and foundation-model approaches has therefore been developed for this task<sup>1-5</sup>, and community benchmarks have begun to standardize their comparison across datasets, cell types and perturbation classes<sup>6,7</sup>. These efforts have established perturbation prediction as a useful abstraction for drug screening, functional genomics and cellular response modeling.

Yet the field still largely operates at the resolution of genes, drugs or experimental conditions. In most current benchmarks, a perturbation is represented as a gene identity, drug identity or dosage condition: the model is asked to predict the effect of knocking down TP53, perturbing KRAS or applying a compound. This design is well matched to gene-knockout and drug-response screens, but it leaves unresolved a finer and clinically central question: how do different protein-coding variants within the same gene alter cellular state? This distinction matters because many disease mechanisms are allele-specific. TP53, for example, carries hundreds of missense mutations with divergent consequences: R175H disrupts protein folding, R273C preserves the folded structure but impairs DNA binding, and R248Q has been associated with dominant-negative and gain-of-function behavior<sup>8</sup>. Similarly, KRAS G12D and G12V are both activating codon-12 mutations, but they differ in downstream signaling behavior and biological effect size<sup>9</sup>. A model that predicts the average consequence of a TP53 perturbation is therefore not necessarily a model that distinguishes TP53 R175H from TP53 R273C. Variant-effect prediction has been extensively benchmarked using multiplexed assays of variant effect and deep mutational scanning, which quantify variant fitness or molecular function at scale<sup>10,11</sup>. However, these resources provide scalar functional scores rather than population-level single-cell transcriptional responses.

Variant-level single-cell measurements and variant-conditioned models have begun to appear, but they do not yet establish a general framework for evaluating held-out protein-coding variant prediction. Single-cell expression-based variant-impact phenotyping has shown that TP53 and KRAS variants can produce reproducible expression signatures when assayed at sufficient depth<sup>12</sup>, and PerturbNet introduced variant-level perturbations into conditional generative modeling in a single-gene GATA1 setting<sup>13</sup>. Recent benchmarking studies have also shown that perturbation-prediction metrics and model rankings can be sensitive to simple baselines, systematic variation and metric geometry<sup>7,14-16</sup>. However, these efforts primarily address gene- or drug-level prediction, or variant phenotyping within measured perturbations. They do not answer whether allele-specific responses within the same gene are distinguishable, predictable or even evaluable under realistic single-cell sampling depth. This evaluation problem is fundamental: standard metrics such as mean squared error or Pearson correlation can reward recovery of the shared direction of a gene-level response even when a model fails to discriminate individual alleles. If split-half variability within a variant approaches the variant-to-wild-type signal, then the measurement itself leaves little signal window for allele-specific ranking.

Here we introduce AllelePerturb, a benchmark and evaluation framework designed specifically for held-out protein-coding variant prediction in single-cell transcriptomics. AllelePerturb-Bench integrates 470 protein-coding variants across four genes (TP53, KRAS, GATA1 and JAK1) and 321,043 single cells (including wild-type controls) spanning three experimental technologies: Perturb-seq, base editing and scSNV-seq. To our knowledge, these datasets span the publicly available single-cell perturbation data that resolve individual protein-coding variants with reusable per-variant labels. AllelePerturb-Eval decomposes prediction quality into three complementary dimensions: direction recovery, allele-level discrimination and differential-expression fidelity. It further defines six generalization splits that test random generalization, positional extrapolation, mechanistic extrapolation, cross-gene transfer, low-depth evaluation and compatibility with existing protocols.

Using this framework, we systematically evaluated 20 feature-model combinations and found a consistent direction-ranking dissociation. All feature-model predictors recovered the transcriptional direction of perturbation effects, with Pearson- $\delta$  values between 0.55 and 0.65, yet none reliably distinguished the responses of held-out variants, with perturbation discrimination scores clustered at the chance expectation of 0.50 (0.49 to 0.52 across predictors, every 95% confidence interval overlapping 0.50). Pub-

lished perturbation models that accept a continuous variant representation (scGen, scVIDR, Biolord and CellFlow), re-scored through an identical harness, reproduced this result, whereas gene-keyed foundation models (GEARS, STATE and scGPT) are allele-blind by construction and a compositional autoencoder did not converge on the continuous inputs; none of the evaluated model classes exceeded the empirical floor, not only the simple baselines. Split-half analysis traced this dissociation to a measurement-resolution floor governed by the ratio between variant-to-wild-type effect size, within-variant sampling noise and effective cells per variant. Extending the same analysis to widely used gene-level Perturb-seq atlases showed that this floor is not unique to allele-level data: at native depth, roughly half of perturbations in the Replogle atlas were un-rankable by direct split-half energy distance. AllelePerturb therefore reframes protein-coding variant prediction from a simple model-ranking problem into a power-aware evaluation problem, and provides a standard framework for determining when variant-level single-cell perturbation prediction is measurable enough to benchmark.

## Results

### AllelePerturb defines perturbation prediction at protein-coding variant resolution

To evaluate perturbation prediction at protein-coding variant resolution, we defined the prediction target as a held-out protein-coding variant rather than a held-out gene, drug or condition. In this setting, the model must predict the transcriptional effect of an unseen allele from its variant features. This differs from gene-level perturbation prediction, in which all perturbations of the same gene are treated as a single condition. In TP53, for example, a gene-level representation collapses variant responses into a common perturbation direction (Fig. 1a, top), whereas allele-resolution analysis preserves variant-specific departures from wild type, with mutations such as R175H, R273C and R248Q occupying distinct regions of the response space (Fig. 1a, bottom). AllelePerturb therefore treats within-gene variant identity as the object of prediction rather than as experimental annotation.

AllelePerturb-Bench integrates four variant-perturbation datasets spanning different proteins, cell types and experimental technologies (Fig. 1b,c). TP53 (393 aa, 98 variants) and KRAS (189 aa, 92 variants) were profiled by Perturb-seq in A549 cells<sup>12</sup>, GATA1 (413 aa, 254 variants) by base editing in hematopoietic stem and progenitor cells<sup>13</sup>, and JAK1 (1,154 aa, 26 variants) by scSNV-seq in HT-29 cells<sup>17</sup>. Together, these datasets comprise 470 protein-coding variants and 321,043 single cells, including the dataset wild-type and control cells. Per-variant cell depth varies substantially across genes, with median depths of 929 and 1,000 cells per variant for TP53 and KRAS, 355 for GATA1 and 104 for JAK1 (Fig. 1c), providing a natural range of sampling regimes for testing variant-level evaluation.

To make held-out variant prediction possible, each variant was represented by protein-level features rather than by a variant identifier. We encoded variants using a 6-dimensional biophysical vector  $\theta$ , including changes in hydrophobicity, side-chain volume and charge, together with fold-core location, functional-switch residue status and hotspot/pathogenic annotation (Fig. 1d). These features describe variant properties beyond gene identity and allow variants from different genes to be placed in a shared feature space. AllelePerturb-Eval then separates prediction quality into three dimensions: direction recovery, allele-level discrimination and differential-expression fidelity. Pearson- $\delta$  measures agreement between predicted and observed perturbation directions, whereas PDS directly tests the defining allele-level requirement: whether a predicted response identifies its target variant among other held-out variants (Fig. 1e). Six split strategies further test random generalization, positional extrapolation, mechanistic extrapolation, cross-gene transfer, low-depth evaluation and compatibility with existing protocols (Fig. 1f), and an end-to-end workflow ties the benchmark together (Fig. 1g). This separation of direction recovery from allele discrimination sets up the central question of the benchmark: whether models can distinguish alleles, rather than merely recover the shared direction of a perturbed gene.

### Variant-level prediction exposes a direction-ranking dissociation

Having defined allele-level discrimination as the central requirement of the task, we next asked whether predictors that recover the direction of a perturbation also identify the correct held-out variant. We eval-



**Figure 1: AllelePerturb defines perturbation prediction at protein-coding variant resolution.** **a**, Conceptual distinction between gene-level and protein-coding variant-level perturbation prediction. Gene-level perturbation identity collapses responses from different variants of the same gene into a single averaged effect, whereas AllelePerturb treats a held-out protein-coding variant as the prediction target and asks whether models can distinguish allele-specific responses within the same gene. **b**, Benchmark coverage: variant positions along the four proteins (TP53, KRAS, GATA1 and JAK1) with annotated domains and hotspot residues; AllelePerturb-Bench integrates 470 protein-coding variants across three experimental technologies (Perturb-seq, base editing and scSNV-seq). **c**, Per-variant cell depth varies substantially across genes, providing a range of sampling regimes for evaluation (98, 92, 254 and 26 variants; 83.4k, 83.6k, 149.2k and 4.9k cells, including wild-type/control, for TP53, KRAS, GATA1 and JAK1, respectively). **d**, Each variant is represented by a 6-dimensional biophysical vector  $\theta$  encoding changes in hydrophobicity, side-chain volume and charge, together with fold-core location, functional-switch residue status and hotspot/pathogenic annotation, placing variants from all genes in a shared feature space. **e**, AllelePerturb-Eval separates prediction quality into direction recovery (Pearson- $\delta$ ), allele-level discrimination (PDS) and differential-expression fidelity. **f**, Six generalization splits test random, positional-extrapolation, mechanistic-extrapolation, cross-gene, low-depth and compatibility settings. **g**, End-to-end benchmark workflow for held-out protein-coding variant prediction.

uated 20 predictors or references, comprising 18 feature-model combinations (6 model heads by 3 feature spaces), the Gene-mean predictor and the WT-null reference (Fig. 2). Across the 19 non-null predictors, allele discrimination remained at chance: PDS-cosine clustered between 0.49 and 0.52 (mean 0.50), and no predictor’s 95% bootstrap confidence interval excluded the 0.50 chance level (Fig. 2b). Point estimates were averaged over 15 independent pseudobulk subsamples to remove single-draw variability, and the interval of every predictor overlapped chance. This conclusion did not depend on the analytic 0.50 baseline: within-split label permutation gave an empirical null mean of 0.500 (range 0.495 to 0.504 across 340 method-split-gene combinations), and only 3.2% of combinations exceeded the permutation 95th percentile, below the 5% false-positive rate expected by chance. PDS-L1 and PDS-L2 gave the same conclusion.

In contrast, the same predictors consistently recovered the transcriptional direction of perturbation effects. Pearson- $\delta$  ranged from 0.55 to 0.65 across predictors, and every non-null predictor had a 95% confidence interval entirely above zero (Fig. 2c). Plotting PDS against Pearson- $\delta$  placed all predictors in the same regime: direction recovered, allele ranking failed (Fig. 2d). Method-to-method variation was small compared with the separation between these two evaluation axes, indicating that variant-level prediction exposes a systematic direction-ranking dissociation rather than a failure of a particular model class.


To make this dissociation quantitative, we decomposed each variant’s effect into a gene-shared component and an allele-specific residual,  $\delta_v = \bar{\delta}_g + r_v$ , and asked what each predictor recovers of the residual. Direction recovery was almost entirely gene-shared: the residual Pearson- $\delta$  was near zero for every predictor (between  $-0.08$  and  $0.03$ ), and a featureless predictor that outputs only the training gene mean already reached the same direction score as the best feature model. Removing the shared component removed the recovered signal, so the models capture the program-level response of a gene rather than the allele-specific residual that distinguishes its variants. Consistent with this, the allele-specific residual was itself unmeasurable for the low-effect genes: an oracle scored on residual profiles remained at chance for TP53 and KRAS (0.48), while it rose for GATA1 and JAK1 (0.59 and 0.89), whose variants carry a resolvable allele-specific signal (Supplementary Table 6).

Differential-expression fidelity provided a biological interpretation of the same dissociation (Fig. 2e). Direction agreement reached 65%, whereas DE-LFC Spearman was 0.25 and DE overlap was only 15%. This gradient (direction 65%, LFC ranking 25%, gene-set overlap 15%) suggests that predictors learn a program-level perturbation direction but progressively lose the finer signal needed to rank effect sizes and identify variant-specific response genes.

The magnitude of the dissociation varied by gene (Fig. 2f). It was strongest for TP53 (0.31), intermediate for KRAS (0.19) and GATA1 (0.13), and near zero for JAK1 ( $-0.04$ ). This ordering anticipated the split-half signal-to-noise analysis below: JAK1 had the strongest per-variant signal and the weakest dissociation, whereas TP53 and KRAS lay near the noise floor. Because each gene here is also a distinct dataset, assay and cell type, this ordering is a per-dataset gradient rather than an intrinsic property of the genes (see Discussion). At the single-variant scale, the same pattern appeared within TP53: Ridge-esm produced stable Pearson- $\delta$  values across representative held-out variants (0.74 to 0.82), but PDS varied from 0.00 to 0.98, with roughly half of examples below chance (Fig. 2g). These results suggest that allele-ranking failure is not explained by model choice alone and motivate a direct test of whether the ground truth itself is measurable at variant resolution.

### Allele discrimination is constrained by an effect-size-to-noise window

The direction-ranking dissociation raised a mechanistic question: do predictors fail to rank alleles because the models are wrong, or because the ground-truth variant responses are not stably measurable at the available depth? To test the measurement component directly, we performed a split-half reproducibility analysis. For each variant, cells were randomly divided into two halves, and we compared the energy distance between the two halves of the same variant ( $D_{\text{self}}$ ) with the distance between that variant and wild type ( $D_{\text{null}}$ ; Fig. 3a,b). A  $D_{\text{self}}/D_{\text{null}}$  ratio near 1 indicates that within-variant sam-



**Figure 2: All methods recover perturbation direction but cannot rank variant identity.** **a**, Unified evaluation of a held-out variant: the predicted response is scored for direction recovery (Pearson- $\delta$ ) and for allele discrimination (PDS, whether it identifies its own variant among the other held-out variants). **b**, Allele-level PDS (cosine) for 20 predictors with 95% bootstrap confidence whiskers; every interval overlaps the chance level of 0.5. **c**, The same predictors recover perturbation direction (Pearson- $\delta$ ), with every interval above zero. **d**, Methods cluster in the direction-without-ranking regime. **e**, Differential-expression fidelity: direction agreement 65%, DE-LFC Spearman 0.25, DE overlap 15%. **f**, The dissociation gap follows a per-dataset gradient across the four genes (TP53 0.31, KRAS 0.19, GATA1 0.13, JAK1 -0.04). **g**, Per-variant scores (Ridge-esm, TP53) show stable direction but unstable ranking.

pling variation is comparable to the variant-to-wild-type signal, leaving little measurement window for allele-specific ranking. A ratio well below 1 indicates that the perturbation signal exceeds the split-half noise and should be rankable by distributional metrics.


The four genes occupied distinct regimes (Fig. 3c). TP53 ( $D_{\text{self}}/D_{\text{null}} = 0.96$ , 95% CI 0.95–0.98), KRAS (1.00, 0.99–1.02) and GATA1 (0.88, 0.86–0.90) lay near the measurement floor, whereas JAK1 was well separated from it (0.21, 0.12–0.33). This ordering matched the dissociation pattern in Fig. 2f: the genes closest to the floor showed the strongest separation between direction recovery and allele discrimination, whereas JAK1 served as a positive control with weak dissociation. GATA1 occupied an intermediate regime, with a long tail of individually rankable variants. These results show that the relevant quantity is not raw cell depth alone but the effect-size-to-noise window available for discrimination. This window is one of identification, not detection: for TP53, KRAS and GATA1 the transcriptional distance between sibling alleles was far smaller than within-variant replicate noise (median nearest-sibling cosine distance 0.23 to 0.46 versus replicate noise near 0.97), so no allele pair was resolvable, whereas the threefold-lower replicate noise of JAK1 left 78% of allele pairs resolvable (Fig. 3g). Detectability against wild type therefore does not imply that alleles can be told apart from one another, which is what discrimination requires. Two independent controls confirmed genuine signal absence rather than a metric artifact. A model-free classifier trained to separate the single cells of two sibling variants under cross-validation reached only chance accuracy for TP53 and KRAS (median AUROC 0.49 and 0.50, against 0.50 under label permutation and a wild-type-versus-wild-type control), rose slightly for GATA1 (0.52) and was high for JAK1 (0.87), so even a supervised model with direct single-cell access cannot separate low-effect sibling alleles (Fig. 3h). The window ratio was moreover invariant across five distance families, from energy and maximum-mean-discrepancy to sliced-Wasserstein, cosine and Euclidean, so the floor is not an artifact of energy-distance geometry (Supplementary Table 4).

This point was clearest when effect size and depth were considered together (Fig. 3d,e). JAK1 had the lowest nominal depth among the four genes (median 104 cells per variant) but the widest effect-size-to-noise window, and it was already saturated in the split-half detection assay at 50 cells per half. In contrast, TP53 and KRAS had much higher nominal depths but narrow allele-specific effect windows, remaining close to the noise floor even with median raw depths of 929 and 1,000 cells per variant. Sub-sampling over 50 to 300 cells per half confirmed the same ordering: JAK1 remained fully detectable across the range, whereas TP53, KRAS and GATA1 increased only modestly and plateaued below 75%. Within the observed depth range, adding cells helps only when the underlying effect window is wide enough for additional sampling to resolve it.

We next asked whether this measurement floor was specific to allele-level perturbation data. Using the same native-depth, all-perturbation, direct split-half energy-distance definition, we applied the rankability analysis to three widely used gene-level Perturb-seq atlases. The three low-effect allele-level genes (TP53, KRAS and GATA1) were 98–100% un-rankable by this definition, whereas JAK1 remained a positive control with only 10% un-rankable variants (Fig. 3f). Among gene-level atlases, Replogle 2022 showed 55.3% un-rankable perturbations (95% CI 52.8–57.5%), compared with 14.6% for Adamson 2016 and 3.4% for Norman 2019 (Fig. 6f). Results were qualitatively robust under alternative rankability criteria, including bootstrap-CI-based, increased-stringency and ratio-threshold definitions (Extended Data Fig. 1). Thus, the detection floor is not a peculiarity of allele-level data: finite-depth distributional distances can leave a substantial fraction of perturbations below the resolution required for stable method ranking.

### Allele-ranking failure persists across evaluation choices and model interfaces

We next tested whether the allele-ranking bottleneck could be explained by evaluation design or variant representation rather than by the measurement regime identified above. Across five generalization splits, Pearson- $\delta$  and PDS behaved differently (Fig. 4a,b). Pearson- $\delta$  remained positive under random, positional-extrapolation and compatibility splits, typically between 0.55 and 0.75, and declined mainly under mechanistic extrapolation and low-depth evaluation. By contrast, PDS remained near the empir-




**Figure 3: Split-half analysis reveals a narrow allele-level measurement window.** **a**, Definition of the measurement window from within-variant ( $D_{\text{self}}$ ) and variant-to-WT ( $D_{\text{null}}$ ) distances. **b**, Representative variants illustrate a closed window (TP53 P222P,  $D_{\text{self}} \approx D_{\text{null}}$ ) and an open window (JAK1 R108Q,  $D_{\text{self}} \ll D_{\text{null}}$ ). **c**, Per-variant  $D_{\text{self}}/D_{\text{null}}$  places TP53 (0.96), KRAS (1.00) and GATA1 (0.88) near the measurement floor and JAK1 (0.21) well below it. **d**, Effect size and effective depth jointly determine the window, not cell number alone. **e**, Depth titration: detection rate rises with cells per half in a gene-dependent manner (JAK1 saturated at 50 cells, others plateau below 75%). **f**, Fraction of variants rankable at native depth (TP53 0%, KRAS 0%, GATA1 2.4%, JAK1 90.0%); gene-level atlases are compared in Fig. 6f. **g**, Allele-allele identification: fraction of sibling-variant pairs whose transcriptional distance exceeds within-variant replicate noise (TP53 and KRAS 0%, GATA1 3.8%, JAK1 78.5%); detection against wild type does not imply identification among alleles. **h**, Classifier two-sample test: cross-validated AUROC separating the single cells of sibling variants is at chance for TP53 and KRAS (0.49, 0.50), marginal for GATA1 (0.52) and high for JAK1 (0.87); label permutation gives 0.50.

ical chance regime across all splits (Fig. 4a,c), with the strongest degradation under the low-depth split. No split produced a consistent above-chance allele-discrimination regime, indicating that the direction-ranking dissociation was not specific to the random holdout setting.

The conclusion was also insensitive to the distance function used for PDS (Fig. 4d). Cosine, L1 and L2 distances yielded nearly identical PDS distributions across the 18 feature-model combinations, all clustered near the chance level of 0.50. Shifting from methods to gene-by-split combinations gave the same picture (Fig. 4f): Pearson- $\delta$  exceeded PDS in 15 of 17 combinations, with the two exceptions occurring in small held-out mechanistic-extrapolation sets where both metrics were close. The GATA1 low-depth split had a PDS of 0.19, consistent with the detection-limit prediction from the split-half analysis.

We then asked whether the bottleneck reflected insufficient variant representation. Under the identical splits, heads and scoring, we compared the 6-dimensional biophysical  $\theta$  vector, 1,280-dimensional ESM-1v embeddings, modern ESM2 embeddings, three mutation-aware ESM2 constructions that encode the change relative to wild type (the mutant-minus-wild-type global embedding, the mutated-residue hidden-state difference and a local-window difference), and AlphaMissense, a purpose-built structure- and evolution-informed variant-effect predictor. Every representation left allele discrimination at chance (PDS 0.46 to 0.48) with similar direction recovery (Pearson- $\delta$  0.63 to 0.65; Fig. 4e, Supplementary Table 7). The direction channel was not feature-driven: a featureless predictor that outputs the training gene mean already reached Pearson- $\delta$  0.67, and shuffling the ESM2 features across variants within a gene left PDS at chance, confirming that the recovered direction is the shared gene-level program rather than allele-specific signal. The failure was also not a matter of the representation discarding the mutation: the mutated-residue difference was highly variant-specific, with within-gene distances 91% of between-gene distances, against 0.1% for the gene-identity-dominated global embedding, yet it too failed. AlphaMissense pathogenicity did not track the measured transcriptional effect size within any gene (Spearman near zero for TP53, KRAS and GATA1). Thus more modern, mutation-aware and structure-aware representations recover the shared perturbation direction but do not overcome the allele-discrimination floor, which lies in the ground-truth measurement rather than in the input.

Finally, we asked whether published perturbation-prediction models fare better than our feature-regression heads under a fair, identical evaluation. We re-scored external models through a single harness: the same held-out variants, pseudobulk profiles, tie-aware PDS and bootstrap and permutation controls used for the in-house heads (Fig. 4g, Table 1; Methods). Four established models that accept a continuous variant representation ran to completion, and all remained at chance: scGen (PDS 0.49, 95% CI 0.46–0.52), scVIDR (0.50, 0.48–0.53), Biolord (0.49, 0.47–0.52) and CellFlow (0.51, 0.48–0.54). None exceeded the feature-model baselines of Fig. 2, whose intervals also overlapped chance. The remaining public tools cannot represent the task: scGPT, GEARS and STATE condition on perturbed-gene identity, so all variants of a gene collapse to a single prediction and only 4 of 470 variants (0.85%) are representable (Fig. 4h). Consistent with this, the gene-constant Gene-mean reference sits at chance (PDS 0.49). A compositional autoencoder (CPA) fitted with a variant-feature encoder did not converge on the continuous inputs. One generative flow, PerturbNet, marginally exceeded the analytic 0.50 chance level (PDS 0.54), but its predictions are confined to a 50-dimensional expression subspace in which random vectors already score 0.52; against a permutation null estimated in that subspace the 0.54-versus-0.52 gap was not significant, and its direction recovery was negligible (Pearson- $\delta$  0.13). We thus tested six model classes: feature regression, latent arithmetic, disentanglement, optimal transport, gene-keyed foundation models and compositional autoencoders. None achieved allele-level discrimination, indicating a measurement floor that bounds every architecture rather than a limitation of any single method.



**Figure 4: Allele-ranking failure persists across evaluation choices and model interfaces.** **a**, PDS across five generalization splits remains at or below chance. **b**, Pearson- $\delta$  stays consistently positive across the same splits. **c**, Empirical permutation null: within-split label permutation centres PDS at 0.50 (per-combination null means 0.495–0.504), with observed exceedances of the 95th percentile at 3.2% (below the expected 5%), so chance-level PDS is not an artifact of the analytic baseline. **d**, Ranking failure is invariant to the distance function (cosine, L1 and L2 all near chance across the 18 feature-model combinations). **e**, Richer features do not close the ranking gap ( $\theta$ , ESM and ESM+ $\theta$  all near 0.50). **f**, The dissociation holds across all gene-by-split combinations, with Pearson- $\delta$  exceeding PDS in 15 of 17. **g**, Published perturbation models scored under one identical harness (same held-out variants, pseudobulk profiles, tie-aware PDS and bootstrap/permutation controls): every variant-conditionable model, the in-house feature-regression range and both references overlap the PDS = 0.50 chance line; PerturbNet’s point estimate (0.54) falls inside its own 50-dimensional subspace null (0.52) (Table 1). **h**, Interface-compatibility audit: gene-keyed models (scGPT, GEARS, STATE) are allele-blind by construction, resolving only 4 of 470 variants (0.85%); this is interface incompatibility, distinct from performance failure.

Table 1: **Published perturbation models under a unified allele-level evaluation.** Perturbation-discrimination score (PDS, cosine) with 95% bootstrap confidence intervals, pooled over five generalization splits and four genes; chance is 0.5. Variant-conditionable models are scored directly; gene-keyed models are allele-blind by construction and reported by resolution coverage. Every method was scored through one harness (identical held-out variants, pseudobulk profiles, tie-aware PDS and permutation null).

| Model family                                                            | Method                  | PDS (95% CI)     | Pearson- $\delta$ |
|-------------------------------------------------------------------------|-------------------------|------------------|-------------------|
| <i>Variant-conditionable (run to completion)</i>                        |                         |                  |                   |
| Optimal-transport flow                                                  | CellFlow                | 0.51 (0.48–0.54) | 0.12              |
| Attribute disentangle-<br>ment                                          | Biolord                 | 0.49 (0.47–0.52) | 0.13              |
| Latent arithmetic (dose)                                                | scVIDR                  | 0.50 (0.48–0.53) | 0.13              |
| Latent arithmetic                                                       | scGen                   | 0.49 (0.46–0.52) | 0.18              |
| Conditional-flow gener-<br>ative                                        | PerturbNet <sup>†</sup> | 0.54 (0.51–0.57) | 0.13              |
| Feature regression                                                      | baselines (Fig. 2)      | 0.49–0.52        | 0.55–0.65         |
| <i>References</i>                                                       |                         |                  |                   |
| Gene identity                                                           | Gene-mean               | 0.49 (0.46–0.52) | 0.66              |
| No perturbation                                                         | WT-null                 | 0.50             | —                 |
| <i>Allele-blind by construction (resolution coverage 4/470 = 0.85%)</i> |                         |                  |                   |
| Foundation model                                                        | scGPT                   | n/a <sup>‡</sup> | —                 |
| Gene knowledge graph                                                    | GEARS                   | n/a <sup>‡</sup> | —                 |
| Virtual-cell (one-hot)                                                  | STATE                   | n/a <sup>‡</sup> | —                 |
| <i>Did not converge</i>                                                 |                         |                  |                   |
| Compositional autoen-<br>coder                                          | variant-CPA             | non-finite loss  | —                 |

<sup>†</sup>PerturbNet generates within a 50-dimensional principal-component subspace in which random predictions already score 0.52; against a permutation null estimated in that subspace its exceedance is not significant. <sup>‡</sup>These models condition on perturbed-gene identity and assign an identical prediction to every variant of a gene, so an allele-level PDS is undefined (a category error).

### The measurement window sets an achievable ceiling and governs benchmark validity

A benchmark can reward a model only for signal that the measurement itself resolves. To bound what is achievable, we computed an oracle predictor whose prediction for each held-out variant is an independent second measurement of that same variant, taken from a disjoint split-half pseudobulk and scored through the identical harness; this defines an empirical replicate-based discrimination ceiling at the specified depth and protocol (Methods). For TP53 and KRAS the ceiling was itself at chance (PDS 0.485, 95% CI 0.445–0.525, and 0.500, 0.440–0.566). This was not an artifact of split-half depth: it was unchanged when all cells were used (0.486 and 0.492), and even a perfect prediction, the deep true mean of a variant, scored against the benchmark’s own noisy measurement remained at chance (0.49 and 0.49; Supplementary Table 8). Allele ranking is therefore unattainable at this resolution by any predictor evaluated against the benchmark truth, and the observed model failure is a measurement limit rather than a modeling one. GATA1 sat marginally above chance (0.572, 0.549–0.595), matching its small resolvable-pair fraction. For JAK1, by contrast, the ceiling reached 0.792 (0.742–0.838) while the best honestly-scored model stayed near chance (0.52), revealing a genuine computational gap in which the measurement is adequate but current models do not exploit it (Fig. 5b,c, Table 2). Consistent with the confounds noted above, the higher JAK1 ceiling reflects its more favorable measurement regime rather than necessarily

stronger allele biology.

The same ceiling governs whether a benchmark can order models at all. We built synthetic predictors of known, graded quality by interpolating between the gene-mean profile and each variant’s measured effect (Fig. 5d,e), and asked how often the observed PDS recovered their true ranking under bootstrap resampling of the held-out variants (Methods). When the oracle window was at chance (ceiling  $\leq 0.52$ ; TP53 and most depths of KRAS), the benchmark recovered the true predictor order only 27% of the time and identified the best predictor only 38% of the time, little better than chance. Once the window opened even modestly (ceiling 0.52–0.56; GATA1), correct ordering rose to 95% and correct winner selection to 97%; above a ceiling of 0.65 both approached 1.0 (Fig. 5g). Recovery also fell when few variants were held out, matching the expected dependence on test-set size. Rankability is therefore not merely descriptive: below the measurement window a benchmark cannot reliably distinguish a good model from a bad one, and any leaderboard it produces is unstable regardless of the models’ true quality (Fig. 5g). This closes the loop from the split-half diagnostic to benchmark validity and motivates predicting rankability before a screen is run. Applying the same resolution analysis to five diverse public perturbation benchmarks placed this floor in context: Norman, Replogle, Adamson and the Virtual Cell Challenge training set all recovered the true predictor ordering (resolution 1.00) and a drug screen (sci-Plex) reached 0.77, whereas TP53 and KRAS remained the only datasets at the floor (0.01 and 0.02; Fig. 5f and Supplementary Table 5). Their unrankability is therefore not a peculiarity of AllelePerturb but a genuinely lower ground-truth resolution than diverse gene- and drug-level screens; consistent with the dependence on test-set size, the many variants of GATA1 (0.51) partly offset its low per-variant signal.

**Table 2: Per-gene measurement window, achievable ceiling and model performance.** Split-half energy-distance ratio  $D_{\text{self}}/D_{\text{null}}$ ; oracle discrimination ceiling (independent second measurement, estimating attainable discrimination); fraction of allele pairs resolvable above replicate noise; fraction of variants rankable at native depth; and best honestly-scored model PDS (canonical multi-seed harness). Chance is 0.5.


| Gene  | $D_{\text{self}}/D_{\text{null}}$ | Oracle ceiling (95% CI) | Pairs resolvable | Rankable (native) | Best model PDS <sup>†</sup> |
|-------|-----------------------------------|-------------------------|------------------|-------------------|-----------------------------|
| TP53  | 0.96                              | 0.485 (0.445–0.525)     | 0%               | 0%                | 0.51                        |
| KRAS  | 1.00                              | 0.500 (0.440–0.566)     | 0%               | 0%                | 0.52                        |
| GATA1 | 0.88                              | 0.572 (0.549–0.595)     | 3.8%             | 2.4%              | 0.54                        |
| JAK1  | 0.21                              | 0.792 (0.742–0.838)     | 78.5%            | 90%               | 0.52                        |

<sup>†</sup>Best model excluding PerturbNet, whose higher per-gene values (0.53–0.54) arise from its 50-dimensional principal-component subspace, in which random predictions already score 0.52 (Table 1).

### Rankability prediction enables power-aware evaluation of variant perturbation screens

If allele discrimination is meaningful only above a measurement floor, then variant-level benchmarks need a way to identify which perturbations are evaluable before full model comparison. We therefore asked whether rankability could be predicted from pre-modeling, pilot-estimable quantities. Using leave-one-dataset-out (LODO) cross-validation across the four AllelePerturb datasets and three external Perturb-seq atlases, we trained a rankability predictor that uses each perturbation’s pilot-estimable pseudobulk effect size as its single feature (Fig. 6a; Methods).

Across the five datasets in which rankability is non-degenerate, effect size separated rankable from un-rankable perturbations with a mean AUROC of 0.96 (Replogle 0.97, Norman 0.93, Adamson 0.93, GATA1 0.95, JAK1 1.00; ROC curves and 95% bootstrap intervals in Fig. 6b,c). TP53 and KRAS could not be evaluated: at native depth every one of their variants is un-rankable, a single-class outcome that is itself the strongest statement that these low-effect genes lie entirely below the measurement floor. Effect size is therefore a usable triage signal exactly where rankability varies; where it does not, the floor has already foreclosed the task. Because this predictor’s feature and label derive from the same cells, we



**Figure 5: The measurement window sets an achievable ceiling and governs benchmark validity.** **a**, The oracle predictor defined: an independent second split-half measurement of each held-out variant, scored through the identical harness, estimating the discrimination attainable at the available depth. **b**, Oracle ceiling per gene versus the best honestly-scored model. TP53 and KRAS ceilings are at chance (measurement-limited, no gap to close); the JAK1 ceiling is high while models remain near chance (a residual computational gap). **c**, Ceiling-versus-model phase map separates two failure modes: TP53 and KRAS are measurement-limited (both near chance), JAK1 is computation-limited (ceiling high, model low), and the upper-right benchmark-solvable regime is where a model would track its ceiling. **d**, Construction of synthetic predictors of known, graded quality by interpolating between the gene-mean profile and each variant's measured effect, fixing the true ordering by design. **e**, Illustration of benchmark validity: the same true model order is scrambled by a near-floor benchmark but recovered by a higher-resolution one. **f**, Benchmark resolution across nine perturbation datasets: the probability of recovering the true ranking of graded synthetic predictors (resolution) versus each dataset's oracle ceiling. The four allele genes and five external gene- and drug-level screens trace a phase transition, collapsing at the floor (TP53, KRAS) and approaching one above a ceiling of  $\sim 0.6$ . **g**, The benchmark-validity transition: full-order recovery and best-model selection probabilities are both low below the measurement window (TP53, KRAS) and jump toward one once the ceiling clears  $\sim 0.52$ – $0.55$  (GATA1) and beyond (JAK1); selecting a winner leads recovering the full order.

also validated it prospectively on disjoint data: from a pilot of only 25 to 50 cells per perturbation we predicted rankability at higher depth on held-out cells, under leave-one-dataset-out cross-validation. Across the three benchmarks with balanced labels (Adamson, Norman and Replogle) the pilot predictor reached AUROC 0.85, 0.94 and 0.99 (mean 0.93), with near-diagonal calibration and a label-permutation null of 0.49; a training-free predictor using the pilot signal-to-noise ratio performed comparably (0.84 to 0.98 per dataset) and the result held at a 25-cell pilot (Fig. 6d). Rankability is therefore prospectively predictable from a small pilot, before a full screen is committed.

The same analysis also reframes experimental depth as an effect-size-conditioned design variable (Fig. 6e). A strong-effect gene such as JAK1 crossed the detection floor from as few as 50 cells per half, whereas TP53 and KRAS remained near the noise floor despite median raw depths of 929 and 1,000 cells per variant. The relevant design target is therefore not a fixed cells-per-variant threshold, but the depth required for a given effect-size, inter-variant-difference and sampling-noise structure. In this sense, the bottleneck shifts from raw cell count to the effect-size-to-noise window that the experiment must clear.


As a practical check, we revisited the external gene-level atlases in an experimental-design context (Fig. 6f). At native depth, 55% of perturbations were un-rankable in Replogle, compared with 15% in Adamson and 3% in Norman. When normalized to the same shallow depth ( $n = 50$  per half), the un-rankable fractions shifted toward one another: Replogle remained high at 55%, while Adamson increased to 39% and Norman to 11%. This pattern indicates that depth normalization can substantially change which perturbations are evaluable, while dataset-specific effect-size structure remains important. Together, these findings motivate a power-aware workflow for variant-level perturbation evaluation (Fig. 6g): estimate effect size and the split-half floor in a pilot assay; triage variants by predicted rankability; size the experiment to clear the effect-size-to-noise floor; and benchmark models only on perturbations whose ground truth is sufficiently measurable. This converts variant-level perturbation evaluation from uniform model ranking into a power-aware, stratified evaluation design.

## Discussion

AllelePerturb shows that protein-coding variant prediction is not simply a finer-grained version of gene-level perturbation prediction. At this resolution, two notions of success that are often coupled in gene-level benchmarks separate: models can recover the shared transcriptional direction of a perturbation while failing to discriminate the responses of individual alleles. This direction-ranking dissociation reframes held-out variant prediction as a measurement-power problem. The central question is not only which model best predicts a perturbation, but whether the allele-specific component of the ground truth is measurable enough to support model comparison.

This finding has implications for both benchmark design and variant biology. Recent studies have shown that perturbation-prediction methods often differ little from simple baselines and that rankings can be sensitive to evaluation choices<sup>14,15</sup>. AllelePerturb provides a complementary metrological explanation: when many perturbations lie below the resolution of distributional distance metrics, model rankings may be unstable regardless of model complexity. Biologically, direction recovery indicates that variants within a gene can share a program-level response, whereas allele discrimination requires resolving the finer signal that distinguishes variants such as TP53 R175H from R273C. These regimes form three nested questions of increasing difficulty, detection (a variant versus wild type), identification (one variant versus another) and prediction (a held-out variant from its features), which map respectively onto assay sensitivity, ground-truth resolution and computational generalization; the failure mode of each dataset locates it on this ladder (Supplementary Table 9). By this account TP53 and KRAS are resolution-limited, JAK1 is model-limited, and only a wider effect-size-to-noise window moves a dataset up the ladder. Our results do not imply that allele-specific transcriptional effects are absent. Prior single-cell variant-impact phenotyping and the JAK1 positive control in AllelePerturb show that such effects can be detectable when the effect-size-to-noise window is sufficiently wide.

The practical consequence is a change in how variant-level perturbation benchmarks should be built. Rather than evaluating all variants uniformly, AllelePerturb-Eval suggests the power-aware workflow



**Figure 6: A power-aware reporting protocol for allele-resolution perturbation prediction.** **a**, Rankability-prediction pipeline: a pilot-estimable effect size triages which perturbations are evaluable before full model comparison, under leave-one-dataset-out cross-validation. **b**, Leave-one-dataset-out ROC curves of the single-feature (log effect-size) rankability predictor. **c**, Dataset-level LODO AUROC (mean 0.96): Replogle 0.97, Norman 0.93, Adamson 0.93, GATA1 0.95, JAK1 1.00; TP53 and KRAS are uniformly un-rankable at native depth (single class), so their AUROC is undefined and they are not evaluable. **d**, Prospective validation: from a 25 to 50-cell pilot, rankability at higher depth is predicted on disjoint held-out cells under leave-one-dataset-out cross-validation (AUROC 0.85, 0.94 and 0.99 for Adamson, Norman and Replogle; label-permutation null 0.49), with near-diagonal calibration. **e**, Effect-size-conditioned design landscape: rankability is set by the effect-size regime, not by a universal cell-count cutoff (filled, rankable; open, un-rankable). **f**, Native versus depth-matched ( $n = 50$  per half) un-rankable fraction on external gene-level atlases: Replogle stays high (55%) while Adamson (15% to 39%) and Norman (3% to 11%) rise toward it, indicating dataset-specific effect-size structure beyond depth alone. **g**, Recommended power-aware allele-resolution workflow.

described above (Fig. 6g), which sizes experiments to the effect-size-to-noise floor and benchmarks models only on perturbations whose ground truth is sufficiently measurable. The LODO rankability predictor implements this idea as a triage tool, not as a universal calculator of absolute cell-number requirements. Several limitations define the current scope: AllelePerturb covers four genes and is enriched for missense variants;  $\theta$  is a deliberately simple biophysical representation; split-half rankability has limited power at very low cell counts; and our audit of external perturbation models reflects public interfaces rather than the theoretical limits of those model families. In addition, gene identity is confounded with assay technology, cell type and stimulation across the four datasets, so the wider measurement window of JAK1 reflects a more favorable effect-size-to-noise regime for that dataset rather than necessarily stronger allele biology.

The four-gene scope reflects the current state of the field as much as a design choice. To our knowledge, AllelePerturb already integrates the coding-variant single-cell perturbation datasets for which allele-resolution transcriptomes and reusable per-variant labels are both publicly available. Additional allele-resolution measurements have since appeared, including SEUSS profiling of RUNX1 variants<sup>18</sup> and transcript-informed single-cell CRISPR sequencing of TP53 and SF3B1<sup>19</sup>. These are valuable resources, but their per-cell variant assignments are not yet deposited in directly reusable form. The raw expression reads are public, whereas the dedicated variant-barcode or genotype-to-cell information required to label individual cells is not archived alongside them. Incorporating such data uniformly would therefore depend on author-provided processed genotypes or bespoke reprocessing, which we did not undertake here. This gap is itself informative. As allele-resolution perturbation screens proliferate, standardized deposition of per-cell variant assignments, not raw reads alone, will become increasingly important. It would let benchmarks such as AllelePerturb expand, and the power-aware evaluation we propose be applied consistently across datasets.

Crossing the detection limit will require progress in both experimental and computational design. Experimentally, larger saturation-variant and base-editing screens with adequate cells per variant, balanced controls and matched replicates will be needed to bring allele-specific ranking into an evaluable regime. Computationally, perturbation models must expose explicit protein-coding variant conditioning rather than treating variants as gene-level perturbation labels. AllelePerturb provides the benchmark, diagnostics and power-aware evaluation framework for this transition, defining when protein-coding variant prediction in single-cell transcriptomics is measurable enough to be meaningfully modeled.

## Methods

### Benchmark construction

**Data sources.** AllelePerturb-Bench integrates single-cell perturbation data for four protein-coding genes assayed across three experimental technologies. TP53 (98 variants) and KRAS (92 variants) were obtained from a Perturb-seq screen of coding variants in A549 lung adenocarcinoma cells (Gene Expression Omnibus GSE161824)<sup>12</sup>. GATA1 (254 variants) was obtained from a base-editing screen in human hematopoietic stem and progenitor cells, using the processed expression matrix and predefined hold-out splits distributed with PerturbNet<sup>13</sup> (Hugging Face repository cyclopet/PerturbNet\_reproduce). JAK1 (26 variants) was obtained from a single-cell single-nucleotide-variant sequencing (scSNV-seq) experiment in HT-29 colorectal cancer cells stimulated with interferon-gamma (European Nucleotide Archive PRJEB48915; processed objects from Zenodo record 10418435)<sup>17</sup>. Cell counts include both variant-assigned cells and the dataset-specific wild-type or control cells retained after preprocessing; the four datasets together comprise 470 protein-coding variants and 321,043 cells (the latter including wild-type and control cells).

**Preprocessing.** For TP53 and KRAS, processed count matrices, cell barcodes and variant-to-cell assignment tables were loaded from GEO, and cells were assigned to variants using the published assignment files. For GATA1, we used the PerturbNet-distributed AnnData object, which provides log-normalized expression over 2,477 highly variable genes together with five predefined holdout partitions. For JAK1, the SingleCellExperiment object was loaded in R (Bioconductor); cells were filtered

to homozygous genotyped calls with unambiguous coding consequences (missense, splice-site, stop-gained or synonymous), the 2,000 most highly variable genes were selected with `scran::modelGeneVar` followed by `getTopHVGs`, and log-normalized counts were exported. Each dataset was processed in its own gene space, because the three cell types do not share a common highly-variable-gene set; the object shared across datasets is the variant feature representation, not the gene space. Within each dataset, every gene was standardized to zero mean and unit variance across all cells, and wild-type cells served as the common reference.

**Variant feature vector.** Each variant  $v$  was represented by protein-level features rather than by a variant identifier, so that held-out variants unseen in training can still be predicted. The primary representation is a six-dimensional biophysical vector

$$\theta_v = (\Delta H_v, \Delta V_v, \Delta Q_v, f_v, s_v, h_v),$$

where the first three components are the mutant-minus-wild-type differences in Kyte-Doolittle hydrophobicity, residue molecular weight and formal charge at pH 7.0. Writing  $p(\cdot)$  for a residue property and  $a_v^{\text{mut}}, a_v^{\text{wt}}$  for the mutant and wild-type amino acids of variant  $v$ , each biophysical difference is  $\Delta p_v = p(a_v^{\text{mut}}) - p(a_v^{\text{wt}})$ , standardized across all 470 variants by  $z(\Delta p_v) = (\Delta p_v - \mu_{\Delta p}) / \sigma_{\Delta p}$ . The remaining three components are structural indicators. The fold-core term  $f_v$  was set to 1.0 for residues within an annotated structural core (the DNA-binding, G-, zinc-finger and kinase domains of TP53, KRAS, GATA1 and JAK1, respectively), to 0.5 for annotated interaction regions, and to 0.0 elsewhere. The functional-switch term  $s_v$  was set to 1.0 for catalytic or metal-coordinating residues. The hotspot or pathogenic marker  $h_v$  was defined exclusively from external prior knowledge (canonical COSMIC and IARC missense hotspot codons, ClinVar pathogenic annotations and gene-specific functional-module membership), with no reference to the single-cell outcome; the full per-gene residue and codon assignments are listed in Supplementary Table 1. An earlier version of  $h_v$  used the top quintile of measured effect size, which leaks the evaluation outcome into the feature; we removed it and re-ran the entire evaluation grid. Ranking performance was unchanged ( $\theta$ -only PDS shifted by at most 0.02 and every method remained at chance), confirming that the central result does not depend on any outcome-derived feature.

**Protein-language-model embeddings.** As an alternative representation, we extracted per-variant embeddings from ESM-1v (`esm1v_t33_650M_UR90S_1`)<sup>20</sup>. Writing  $\mathbf{H}_v \in \mathbb{R}^{L_v \times 1280}$  for the final-layer hidden states of the full-length mutant sequence of length  $L_v$ , the embedding is the mean over residue positions,

$$e_v = \frac{1}{L_v} \sum_{\ell=1}^{L_v} \mathbf{H}_v[\ell, :] \in \mathbb{R}^{1280},$$

used directly, without dimensionality reduction, as the ESM feature space.

### Evaluation protocol (AllelePerturb-Eval)

**Pseudobulk perturbation profiles.** Let  $C_v$  and  $C_{\text{WT}}$  denote the cells of variant  $v$  and of wild type, and  $x_i \in \mathbb{R}^G$  the standardized expression of cell  $i$  over  $G$  genes. The perturbation profile of variant  $v$  is the pseudobulk difference

$$\delta_v = \frac{1}{|C_v|} \sum_{i \in C_v} x_i - \frac{1}{|C_{\text{WT}}|} \sum_{j \in C_{\text{WT}}} x_j.$$

To bound the influence of the most deeply sampled variants, up to 300 cells were sampled without replacement per variant and per wild-type reference; variants with fewer than five cells were excluded. Every predictor maps a variant feature vector to an estimate  $\hat{\delta}_v$  of this profile, and all metrics compare  $\hat{\delta}_v$  to the measured  $\delta_v$ .

AllelePerturb-Eval separates prediction quality into three families of metrics. **Direction recovery** asks whether a predicted profile points the right way in expression space. **Allele discrimination** asks whether a predicted profile is specific enough to identify its own variant among the other held-out variants. **Differential-expression fidelity** asks whether the predicted and measured profiles agree on which genes change and by how much. The ten metrics are defined as follows.

**Perturbation discrimination score (PDS).** PDS is the primary metric and directly tests allele discrimination: is a predicted profile closer to its own measured profile than to those of the other held-out variants? This nearest-match discrimination logic follows the perturbation-discrimination principle used in recent virtual-cell evaluation efforts<sup>15,16</sup>, adapted here to protein-coding variants and computed with tie-aware ranking. Let  $T$  be the set of  $n_{\text{test}} = |T|$  held-out variants and  $d(\cdot, \cdot)$  a distance. For variant  $v$ , the rank of the correct match among all candidates is

$$r_v = |\{ w \in T : d(\hat{\delta}_v, \delta_w) < d(\hat{\delta}_v, \delta_v) \}| + 1,$$

with tied distances assigned the average rank, and the score is normalized so that the nearest match scores 1 and random matching scores 0.5 in expectation,

$$\text{PDS}(v) = 1 - \frac{r_v - 1}{n_{\text{test}} - 1}.$$

PDS was computed under three distances,

$$d_{\cos}(a, b) = 1 - \frac{a^\top b}{\|a\| \|b\|}, \quad d_1(a, b) = \|a - b\|_1, \quad d_2(a, b) = \|a - b\|_2,$$

giving PDS\_cos (primary), PDS\_L1 and PDS\_L2. The WT-null reference predicts an identical zero profile for every variant, so all of its distances tie; average-rank assignment places it at exactly 0.5 under all three distances (the cosine distance to a zero vector is undefined and was treated as a uniform tie). Because the empirical chance level under small, gene-stratified test sets with ties can depart slightly from the analytic 0.5, we additionally estimated it by label permutation (see Statistics).

**Direction-recovery metrics.** Averaging the Pearson correlation over held-out variants gives

$$\text{Pearson-}\delta = \frac{1}{n_{\text{test}}} \sum_{v \in T} \text{corr}(\hat{\delta}_v, \delta_v),$$

and  $\delta$ -cosine is the analogous average of  $\hat{\delta}_v^\top \delta_v / (\|\hat{\delta}_v\| \|\delta_v\|)$ . Pearson- $\delta$  (top-20) restricts the correlation to the 20 genes with the highest variance across the training-set profiles, testing whether agreement is concentrated in a small number of high-variance genes.

**Differential-expression metrics.** For each variant, differentially expressed genes were defined by a per-gene Welch  $t$ -test of variant against wild-type cells (up to 300 cells per group). Let  $\mathcal{D}_v^{\text{true}}$  be the  $k = 50$  genes with the largest absolute  $t$ -statistic and  $\mathcal{D}_v^{\text{pred}}$  the  $k$  genes with the largest absolute predicted change. DE overlap is

$$\text{DE overlap}(v) = \frac{|\mathcal{D}_v^{\text{pred}} \cap \mathcal{D}_v^{\text{true}}|}{k}.$$

DE-LFC Spearman is the Spearman correlation between measured log-fold-changes and predicted changes over  $\mathcal{D}_v^{\text{true}}$ , and direction agreement is the fraction of genes in  $\mathcal{D}_v^{\text{true}}$  for which  $\text{sign}(\hat{\delta}_{v,g}) = \text{sign}(\delta_{v,g})$ .

**Reconstruction metric.** Mean absolute error is  $\text{MAE}(v) = G^{-1} \sum_g |\hat{\delta}_{v,g} - \delta_{v,g}|$ , averaged over held-out variants.

## Generalization splits

Six split strategies probe distinct axes of generalization. In each, models are trained on the training variants of a gene and evaluated on its held-out variants.

1. Random: 35% of variants per gene held out uniformly at random, using the predefined partition where one is available (GATA1).
2. Positional (out-of-distribution position): variants in the C-terminal half of each protein held out.
3. Mechanistic (out-of-distribution mechanism): variants at hotspot or functional-switch residues held out.
4. Cross-gene: all GATA1 variants held out with training on the remaining genes. This split is reported separately and excluded from the main grid, because the datasets do not share a gene space and cross-gene transfer is therefore not defined for a shared decoder.
5. Low-depth: variants with fewer than 200 cells held out, isolating the shallow-sampling regime.
6. Compatibility: held-out sets matched to the PerturbNet evaluation protocol for GATA1, enabling comparison with prior work.

## Prediction methods

We evaluated 20 predictors: 18 feature-model combinations formed by crossing six regression heads with three feature spaces, plus two references. Each predictor learns a map  $g_\phi : \mathbb{R}^p \rightarrow \mathbb{R}^G$  from a  $p$ -dimensional variant feature vector to a  $G$ -dimensional pseudobulk profile, fit by minimizing squared error on the training variants,  $\min_\phi \sum_{v \in \text{train}} \|g_\phi(\text{feat}_v) - \delta_v\|_2^2$  (with the head-specific regularizer below), and scored by the identical protocol above.

**Regression heads.** We used six standard regression heads spanning linear, tree-ensemble, neighbour and neural-network families: ridge, lasso, random forest, gradient-boosted trees,  $k$ -nearest-neighbours regression and a multilayer perceptron (hyperparameters in Supplementary Table 2). Linear and neighbour heads regress directly onto the full expression profile. Because tree ensembles scale poorly to thousands of output genes, for output spaces larger than 200 genes the random forest and gradient-boosted heads were fit on a  $q$ -component PCA of the training profiles, with predictions mapped back to gene space by the inverse transform. All heads used a fixed random seed.

**Feature spaces.**  $\theta$  (six-dimensional biophysical vector); ESM (1,280-dimensional ESM-1v embedding); and ESM+ $\theta$  (their 1,286-dimensional concatenation). As a representation robustness control (Supplementary Table 7) we additionally extracted ESM2-650M final-layer representations of the mutant sequence, three mutation-aware constructions of the mutant-minus-wild-type difference (the global mean difference, the mutated-residue hidden-state difference  $h_p^{\text{mut}} - h_p^{\text{wt}}$ , and a  $\pm 16$ -residue local-window difference; proteins longer than 1,022 residues were cropped to a window centred on the mutation, and multi-substitution variants averaged their site differences), and AlphaMissense pathogenicity scores; all were scored through the identical splits, heads and metrics.

**References.** The WT-null reference predicts  $\hat{\delta}_v = 0$  for every variant, quantifying the score attainable with no perturbation signal. The Gene-mean reference predicts the mean training-variant profile of the held-out variant's own gene,  $\hat{\delta}_v = |\text{train}(g_v)|^{-1} \sum_{w \in \text{train}(g_v)} \delta_w$ ; because it uses the gene label  $g_v$ , it has access to information a variant-blind predictor does not, and we report it as an informative upper reference on how much of the response is explained by gene identity alone rather than as a competing method. Gene-mean is defined only for genes with at least one training variant in a split and is therefore omitted under the cross-gene split.

**Published models under a unified evaluation.** To compare against published perturbation-prediction methods on identical footing, every model was scored through one harness: the same expression arrays, deterministic pseudobulk profiles, tie-aware PDS and bootstrap and permutation procedures used for

the in-house heads, across the five main splits (Table 1). Four models that accept a continuous variant representation were trained per gene on the training variants and their wild-type control, and predicted each held-out variant from its feature vector: scGen and scVIDR<sup>1,21</sup> (latent vector arithmetic,  $\theta$  as the perturbation attribute), Biolord<sup>22</sup> (ordered-attribute disentanglement) and CellFlow<sup>23</sup> (conditional optimal-transport flow, with each variant’s ESM-1v embedding as the condition, including held-out variants so unseen alleles resolve by feature). We additionally adapted CPA<sup>2</sup> in a chemCPA style, replacing its categorical perturbation dictionary with a frozen encoding of the variant features; this model produced non-finite training losses on the continuous inputs and did not converge, even with gradient clipping and reduced learning rates. Three further models (scGPT<sup>5</sup>, GEARS<sup>4</sup> and STATE<sup>24</sup>) condition on perturbed-gene identity and assign an identical prediction to every variant of a gene; we therefore report their allele-resolution coverage (4 of 470 variants) rather than a discrimination score, which is undefined for a gene-invariant predictor. PerturbNet<sup>13</sup>, a conditional-flow generative model conditioned on ESM-1v, generates within a 50-dimensional principal-component subspace; because random predictions confined to this subspace already score 0.52, we evaluated it against a permutation null estimated in the same subspace. Confidence intervals are 2,000-sample bootstraps over held-out variants; permutation nulls use 1,000 prediction-to-target shuffles, as for the in-house grid.

### Split-half rankability analysis

To separate model limitation from measurement limitation, we asked whether variant responses are stably measurable at the available sampling depth. The energy distance between two cell sets  $X = \{x_i\}_{i=1}^n$  and  $Y = \{y_j\}_{j=1}^m$  is

$$E(X, Y) = \frac{2}{nm} \sum_{i=1}^n \sum_{j=1}^m \|x_i - y_j\| - \frac{1}{n^2} \sum_{i,i'} \|x_i - x_{i'}\| - \frac{1}{m^2} \sum_{j,j'} \|y_j - y_{j'}\|,$$

with Euclidean norms. For each variant with at least 50 cells, cells were partitioned into two random halves  $A$  and  $B$ , and a size-matched wild-type control sample  $C$  was drawn. We computed the within-variant replicate distance  $D_{\text{self}} = E(A, B)$  and the variant-to-wild-type distance  $D_{\text{null}} = E(A, C)$ . Distances were evaluated over a range of subsample depths per half and in three representation spaces, with 50 random seeds per configuration; the canonical values use a 50-component PCA at native maximum depth (Supplementary Table 3). Each seed distribution was summarized by its median and its 2.5 to 97.5 percentile interval.

A variant was classified as rankable when its signal exceeded the sampling width of the replicate distance,

$$S = \text{med}(D_{\text{null}}) - \text{med}(D_{\text{self}}) > W = Q_{97.5}(D_{\text{self}}) - Q_{2.5}(D_{\text{self}}),$$

and un-rankable otherwise. This signal-versus-noise criterion is the primary definition; we confirmed that the conclusions are qualitatively unchanged under four alternatives (lower confidence bound of  $S$  above zero;  $S > 1.5W$ ;  $S > 2W$ ; and a ratio-based threshold), which shift absolute fractions but preserve the ordering across datasets (Extended Data Fig. 1). The same analysis was applied to the four AllelePerturb genes and to three widely used gene-level Perturb-seq atlases, Replogle 2022, Norman 2019 and Adamson 2016 (K562 cells), obtained from the scPerturb data portal<sup>7</sup>, using an identical native-depth, all-perturbation, direct split-half definition throughout. Here native depth denotes each perturbation evaluated at the deepest split-half subsample rung it supports (from 25 to 800 cells per half) in the 50-component PCA space, so that shallower datasets are scored at correspondingly shallower depth; the resulting per-dataset un-rankable fractions are those reported in Fig. 3f and Fig. 6f.

### Rankability predictor

To test whether rankability can be anticipated before a full model comparison, we trained a logistic-regression classifier to predict per-perturbation rankability under leave-one-dataset-out cross-validation. Each perturbation contributed a single example: its canonical native-depth rankability label (deepest

available split-half bin, energy distance, PCA space, as above) and its pilot-estimable pseudobulk effect size  $\|\delta_v\|_2$ . We excluded the split-half quantities that define the label ( $S$ ,  $W$ , their ratio,  $D_{\text{self}}$ ,  $D_{\text{null}}$ ) and any evaluation-configuration identity (distance metric, representation space, subsample depth), which are not properties of the perturbation and would leak the label. Discrimination was measured by the area under the receiver-operating-characteristic curve with 2,000-sample bootstrap intervals over held-out perturbations; held-out datasets whose label was single-class (all un-rankable, namely TP53 and KRAS at native depth) are reported as not evaluable rather than assigned an undefined AUROC. The predictor is implemented in `scripts/figures/rankability_predictor.py`.

**Prospective pilot-to-full validation.** To test that the rankability predictor is prospective rather than a retrospective consequence of same-source data, we re-evaluated it on cells disjoint from those defining the label. For each perturbation a single random permutation assigned the first 50 cells to a pilot set and the next  $2T$  cells to an evaluation set; the pilot effect size  $\|\hat{\delta}\|_2$ , and a pilot signal-to-noise ratio from a pilot split-half, were computed from 25 or 50 pilot cells, while the rankability label at target depth  $T \in \{100, 200\}$  was computed by the split-half criterion on the disjoint evaluation cells. A leave-one-dataset-out logistic predictor and a training-free pilot signal-to-noise threshold were evaluated across the seven datasets; because pilot and evaluation cells never overlap, this is a genuine prospective test. We report per-dataset AUROC over the datasets with balanced labels rather than a pooled value, which is inflated by cross-dataset base-rate differences, and a within-dataset label-permutation control returned a null AUROC of 0.49. Scripts are provided under `results/pilot_validation/`.

### Measurement ceiling, pairwise identification and benchmark resolution

Three analyses separate measurement limits from model limits; all reuse the identical deterministic pseudobulk and tie-aware PDS of the evaluation protocol, average over 15 random cell subsamples, and require no additional data.

**Oracle ceiling.** For each variant with at least five cells, cells were partitioned into two disjoint halves and a size-matched wild-type control was split likewise, giving two independent pseudobulk deltas. Using one half as the prediction and the other as the evaluation truth, we scored PDS through the identical harness (same candidate sets, tie-aware ranking and split structure as the model comparison). Because the prediction is a genuine second measurement of the same variant, its PDS defines an empirical replicate-based discrimination ceiling at that depth and protocol; we report it as such rather than as a mathematical upper bound. PerturbNet is excluded from this per-gene comparison because it is scored in a reduced principal-component subspace that is not comparable to the full-gene-space oracle.

**Ceiling sensitivity and gene-shared decomposition.** We checked that the ceiling is a genuine measurement limit rather than an artifact of the split-half protocol. Recomputing the oracle using all available cells per half rather than the 300-cell cap left it unchanged, and substituting a near-true prediction (the deep mean of a variant from cells disjoint from the evaluation truth) scored against the benchmark’s noisy pseudobulk truth also left it at chance for the low-effect genes, indicating that the limit lies in the ground-truth resolution rather than in prediction quality; shrinking predictions toward the gene mean only lowered discrimination, so mean-shrinkage denoising does not raise the ceiling. Separately, we decomposed each effect into a gene-shared mean and an allele-specific residual,  $\delta_v = \bar{\delta}_g + r_v$ , and re-computed the oracle and the per-predictor Pearson- $\delta$  on the residual profiles to quantify how much of direction recovery is gene-shared.

**Pairwise identification.** In the same cosine-delta space, we computed the split-half replicate distance  $D_{\text{self}}(v)$  and the pairwise distance  $D(v_i, v_j)$  between variant mean deltas at a matched half-depth. A pair was called resolvable when  $D(v_i, v_j)$  exceeded the mean of the two variants’ replicate distances, and a variant identifiable when its nearest sibling lay beyond its own replicate noise.

**Metric-family and classifier controls.** To show the floor is not specific to energy distance, we recomputed the split-half window ratio  $D_{\text{self}}/D_{\text{null}}$  in the 50-component PCA space under energy distance, maximum-mean-discrepancy (RBF kernel, median-heuristic bandwidth), sliced-Wasserstein, cosine and Euclidean distances (Supplementary Table 4). As a distribution-free identification test, we trained an

$\ell_2$ -regularized logistic-regression classifier to separate the single cells of two conditions under five-fold cross-validation in the same space, with the principal components fit label-blind and classes balanced by subsampling; we report the area under the ROC curve for variant-versus-variant (identification) and variant-versus-wild-type (detection) comparisons. Because all variants of a gene derive from one pooled experiment, there is no per-variant batch confound; label-permutation and wild-type-versus-wild-type controls both returned a chance AUROC of 0.5, and where the classifier exceeded chance its accuracy survived regressing out per-cell depth.

**Benchmark resolution.** To test whether the measurement window governs benchmark validity, we built synthetic predictors of known, graded quality,  $\hat{\delta}(v, \alpha) = (1 - \alpha) \bar{\delta}_g + \alpha \delta_v^{\text{build}}$ , where  $\delta_v^{\text{build}}$  is a pseudobulk delta from a build subsample disjoint from the evaluation truth and  $\bar{\delta}_g$  is the gene mean; true quality increases monotonically with  $\alpha$ , and at  $\alpha = 1$  the attainable score equals the oracle ceiling. Scoring the family through the harness and bootstrapping the held-out variants, we recorded the probability of recovering the full  $\alpha$ -ordering and of selecting the top predictor, as a function of the oracle window and of the number of held-out variants.

**External benchmark calibration.** We applied the identical controlled-predictor resolution analysis to five public perturbation datasets (Norman, Replogle, Adamson, the Virtual Cell Challenge training set and sci-Plex), loading each one’s ground-truth single-cell profiles and computing the resolution, the probability of recovering the true predictor ordering under bootstrap over perturbations, at 50 cells per half over perturbations with at least 100 cells (Supplementary Table 5). Because this inclusion threshold and depth differ from the native-depth main analysis, the allele perturbation counts and rankable fractions there differ from Table 2. Analysis scripts are provided under `results/benchmark_resolution/`.

**Analytical scaling.** Under a finite-sample pseudobulk model, two half-profiles of a variant satisfy  $\mathbb{E}\|d_A - d_B\|^2 = 2\eta^2$  with per-profile noise  $\eta^2 = \text{tr}(\Sigma)/m$ , and two variants satisfy  $\mathbb{E}\|d_i - d_j\|^2 = \|\mu_i - \mu_j\|^2 + 2\eta^2$ , giving  $(D_{\text{self}}/D_{\text{null}})^2 = 1/(1 + S)$  with  $S = \|\delta\|^2/(2\eta^2)$  and a dimensionless discrimination signal-to-noise  $\rho = \bar{\Delta}\sqrt{m}/(\sqrt{2}\eta)$ ; discrimination depends on the data only through  $\rho$  and the number of candidate variants. We confirmed the derived depth scaling empirically ( $\eta^2 \propto 1/m$  across all genes). As a direct statement of the floor, the noise-debiased between-variant signal  $\bar{\Delta}^2$  lay at or below the sampling-noise resolution for TP53, KRAS and GATA1 at every tested depth, so only JAK1 traced the increasing branch of  $\rho$ . The scaling is derived; the functional form is calibrated from the synthetic sweep rather than assumed for the tie-aware cosine score.

## Power analysis

To relate detection to sampling depth, cells were subsampled at 50, 100, 150 and 300 per variant and the split-half rankability test was recomputed at each depth. The fraction of variants classified as detectable was plotted against subsample size for each gene, and a detection threshold (the depth at 50% detection) was estimated by sigmoid fitting where the data supported it.

## Statistics

Confidence intervals for PDS and Pearson- $\delta$  were obtained by nonparametric bootstrap over held-out variants (2,000 resamples, percentile method, 95% level). A method’s PDS was called to differ from chance when its 95% confidence interval excluded 0.50, and its direction significantly positive when the lower bound of its Pearson- $\delta$  interval exceeded 0. The per-gene  $D_{\text{self}}/D_{\text{null}}$  ratio (Fig. 3c) and the external-atlas rankability fractions (Fig. 6f) were bootstrapped over variants and over perturbations, respectively. To confirm that the analytic PDS chance level of 0.50 holds under the actual tie structure and small test sets, we ran a within-split label-permutation control: for each method-split-gene combination the full prediction-to-truth distance matrix was precomputed and the prediction-to-target correspondence was permuted 1,000 times, recomputing PDS with the same tie-aware ranking each time. The empirical null mean was 0.500 (range 0.495 to 0.504 across 340 combinations), and only 3.2% of combinations exceeded the permutation 95th percentile, consistent with the 5% false-positive rate expected by chance. All bootstraps and permutations used fixed random seeds. We report intervals to charac-

terize uncertainty rather than as per-method hypothesis tests: the conclusions rest on the direction and consistency of the effect across all 19 non-reference predictors, which move together, and so do not depend on any single comparison or on a multiplicity correction. To ensure that discrimination estimates did not depend on a single stochastic pseudobulk draw, PDS point estimates in the method comparison were averaged over 15 independent random subsamples of cells per variant, with bootstrap confidence intervals computed on the subsample-averaged per-variant scores.

### Software and reproducibility

The evaluation grid was run in Python 3.11 with NumPy 2.4.6, pandas 2.3.3, scikit-learn 1.7.2 and SciPy 1.14.1. Single-cell loading and preprocessing used scanpy 1.10 (Python) and scan with SingleCellExperiment (R, Bioconductor). ESM-1v embeddings were extracted with fair-esm 2.0.0 and PyTorch 2.10. Figures were rendered with matplotlib 3.10. All code, processed benchmark tables and evaluation protocols are available at <https://github.com/Boom5426/AllelePerturb>.

### Data availability

All single-cell datasets analysed in this study are publicly available. The TP53 and KRAS Perturb-seq data are available from the Gene Expression Omnibus under accession GSE161824<sup>12</sup>. The GATA1 base-editing data were originally generated by Perturb(BE)-seq in human hematopoietic stem and progenitor cells and deposited at the Gene Expression Omnibus under accession GSE215253; we used the processed expression matrix and predefined holdout splits redistributed with PerturbNet (Hugging Face repository [cyclopet/PerturbNet\\_reproduce](https://huggingface.co/cyclopet/PerturbNet_reproduce))<sup>13</sup>. The JAK1 scSNV-seq data are available from the European Nucleotide Archive under accession PRJEB48915, with processed objects at Zenodo (<https://doi.org/10.5281/zenodo.10418435>)<sup>17</sup>. The gene-level Perturb-seq atlases used for the rankability analysis (Replogle 2022, Norman 2019 and Adamson 2016) were obtained from the scPerturb data portal<sup>7</sup>. The derived AllelePerturb-Bench tables that support all main and Extended Data figures (per-variant pseudobulk profiles, variant feature vectors, generalization-split definitions and the full evaluation grids) are available at <https://github.com/Boom5426/AllelePerturb>.

### References

- [1] Mohammad Lotfollahi, F. Alexander Wolf, and Fabian J. Theis. scGen predicts single-cell perturbation responses. *Nature Methods*, 16:715–721, 2019. doi: 10.1038/s41592-019-0494-8.
- [2] Mohammad Lotfollahi, Anna Klimovskaia Susmelj, Carlo De Donno, et al. Predicting cellular responses to complex perturbations in high-throughput screens. *Molecular Systems Biology*, 19:e11517, 2023. doi: 10.15252/msb.202211517.
- [3] Charlotte Bunne, Stefan G. Stark, Gabriele Gut, et al. Learning single-cell perturbation responses using neural optimal transport. *Nature Methods*, 20:1759–1768, 2023. doi: 10.1038/s41592-023-01969-x.
- [4] Yusuf Roohani, Kexin Huang, and Jure Leskovec. Predicting transcriptional outcomes of novel multigene perturbations with GEARS. *Nature Biotechnology*, 42:927–935, 2024. doi: 10.1038/s41587-023-01905-6.
- [5] Haotian Cui, Chloe Wang, Hassaan Maan, et al. scGPT: toward building a foundation model for single-cell multi-omics using generative AI. *Nature Methods*, 21:1470–1480, 2024. doi: 10.1038/s41592-024-02201-0.
- [6] Thomas M. Norman, Max A. Horlbeck, Joseph M. Replogle, et al. Exploring genetic interaction manifolds constructed from rich single-cell phenotypes. *Science*, 365:786–793, 2019. doi: 10.1126/science.aax4438.
- [7] Stefan Peidli, Tessa D. Green, Ciyue Shen, et al. scPerturb: harmonized single-cell perturbation data. *Nature Methods*, 21:531–540, 2024. doi: 10.1038/s41592-023-02144-y.

- [8] Evan H. Baugh, Hua Ke, Arnold J. Levine, Richard A. Bonneau, and Chris Sander Chan. Why are there hotspot mutations in the TP53 gene in human cancers? *Cell Death & Differentiation*, 25:154–160, 2018. doi: 10.1038/cdd.2017.180.
- [9] John C. Hunter, Arianna Manandhar, Mario A. Carrasco, et al. Biochemical and structural analysis of common cancer-associated KRAS mutations. *Molecular Cancer Research*, 13:1325–1335, 2015. doi: 10.1158/1541-7786.MCR-15-0203.
- [10] Pascal Notin, Aaron W. Kollasch, Daniel Ritter, et al. ProteinGym: large-scale benchmarks for protein fitness prediction and design. In *Advances in Neural Information Processing Systems*, volume 36, pages 64331–64379, 2023.
- [11] Daniel Esposito, Jochen Weile, Jay Shendure, et al. MaveDB: an open-source platform to distribute and interpret data from multiplexed assays of variant effect. *Genome Biology*, 20:223, 2019. doi: 10.1186/s13059-019-1845-6.
- [12] Oana Ursu, James T. Neal, Dhawal Bhatt, et al. Massively parallel phenotyping of coding variants in cancer with Perturb-seq. *Nature Biotechnology*, 40:896–905, 2022. doi: 10.1038/s41587-021-01160-7.
- [13] Zhiyuan Yu and Joshua D. Welch. PerturbNet predicts single-cell responses to unseen chemical and genetic perturbations. *Molecular Systems Biology*, 21:960–982, 2025. doi: 10.1038/s44320-025-00131-3.
- [14] Constantin Ahlmann-Eltze, Wolfgang Huber, and Simon Anders. Deep-learning-based gene perturbation effect prediction does not yet outperform simple linear baselines. *Nature Methods*, 22:1657–1661, 2025. doi: 10.1038/s41592-025-02772-6.
- [15] Rocío Viñas Torné, Maciej Wiatrak, Zoe Piran, et al. Systema: a framework for evaluating genetic perturbation response prediction beyond systematic variation. *Nature Biotechnology*, 44:1050–1059, 2026. doi: 10.1038/s41587-025-02777-8.
- [16] Charlotte Bunne, Yusuf Roohani, Yanay Wenteler, et al. How to build the virtual cell with artificial intelligence: priorities and opportunities. *Cell*, 187:7045–7063, 2024. doi: 10.1016/j.cell.2024.11.015.
- [17] Sarah E. Cooper, Matthew A. Coelho, Magdalena E. Strauss, et al. scSNV-seq: high-throughput phenotyping of single nucleotide variants by coupled single-cell genotyping and transcriptomics. *Genome Biology*, 25:20, 2024. doi: 10.1186/s13059-024-03169-y.
- [18] Kivilcim Ozturk, Rebecca Panwala, Jeanna Sheen, Kyle Ford, Nathan Jayne, Andrew Portell, Dong-Er Zhang, Stephan Hutter, Torsten Haferlach, Trey Ideker, Prashant Mali, and Hannah Carter. Interface-guided phenotyping of coding variants in the transcription factor RUNX1. *Cell Reports*, 43:114436, 2024. doi: 10.1016/j.celrep.2024.114436.
- [19] Heon Seok Kim, Susan M. Grimes, Tianqi Chen, Anuja Sathe, Billy T. Lau, Gue-Ho Hwang, Sangsu Bae, and Hanlee P. Ji. Direct measurement of engineered cancer mutations and their transcriptional phenotypes in single cells. *Nature Biotechnology*, 42:1254–1262, 2024. doi: 10.1038/s41587-023-01949-8.
- [20] Joshua Meier, Roshan Rao, Robert Aber, et al. Language models enable zero-shot prediction of the effects of mutations on protein function. In *Advances in Neural Information Processing Systems*, volume 34, pages 29287–29303, 2021.
- [21] Omar Kana, Rance Nault, David Filipovic, Daniel Marri, Tim Zacharewski, and Sudin Bhattacharya. Generative modeling of single-cell gene expression for dose-dependent chemical perturbations. *Patterns*, 4(8):100817, 2023. doi: 10.1016/j.patter.2023.100817.

- [22] Zoe Piran, Niv Cohen, Yedid Hoshen, and Mor Nitzan. Disentanglement of single-cell data with biolord. *Nature Biotechnology*, 42:1678–1683, 2024. doi: 10.1038/s41587-023-02079-x.
- [23] Dominik Klein, Jonas Simon Fleck, Daniil Bobrovskiy, et al. CellFlow enables generative single-cell phenotype modeling with flow matching. *bioRxiv*, 2025. doi: 10.1101/2025.04.11.648220.
- [24] Abhinav Adduri, Dhruv Gautam, et al. Predicting cellular responses to perturbation across diverse contexts with state. *bioRxiv*, 2025. doi: 10.1101/2025.06.26.661135.