

Figure 5 (placeholder — to be drawn)

The measurement window sets an achievable ceiling and governs benchmark validity.

- (a) Oracle discrimination ceiling per gene vs best honestly-scored model (TP53 0.485, KRAS 0.500 at chance; GATA1 0.572; JAK1 0.792 while models near 0.52).
- (b) Benchmark-resolution curve: $P(\text{recover true model order})$ and $P(\text{select best})$ vs oracle window — 0.27/0.38 at ceiling ≤ 0.52 , rising to ~ 1.0 above 0.65.